

KUBERNETES

Kubernetes (K8s) is an open-source platform to manage containerized applications across clusters of nodes. It automates deployment, scaling, and operations.

Features:

- **Container Orchestration:** Kubernetes manages the lifecycle of containers, including their deployment, scaling, and updates.
- **Scalability and Resilience:** Kubernetes ensures that applications can scale up or down based on demand and can automatically recover from failures, ensuring high availability.
- **Service Discovery:** Kubernetes provides mechanisms for services to discover each other, making it easy to connect different parts of a distributed application.
- **Configuration Management:** Kubernetes offers tools like Secrets and ConfigMaps to securely manage application configurations and sensitive data.
- **Automation:** Kubernetes automates many of the manual tasks involved in managing containerized applications, reducing operational overhead.

Core Concepts

1. Cluster

- A **Kubernetes cluster** is a set of nodes grouped together to run containerized applications.
- It consists of a **control plane (master node)** and one or more **worker nodes**.

2. Node

- A **node** is a virtual or physical machine in the cluster.
- Each node runs pods and is managed by the control plane.
- Two types:

- **Master Node**
- **Worker Node**

3. Master Node (Control Plane)

- Manages the cluster.
- Responsible for:
 - Scheduling pods
 - Maintaining cluster state
 - Monitoring and responding to events
- **Key components:**
 - [kube-apiserver](#): Frontend of the control plane
 - [etcd](#): Key-value store for cluster data
 - [kube-scheduler](#): Assigns pods to nodes
 - [kube-controller-manager](#): Runs controllers (like node controller, replication controller)

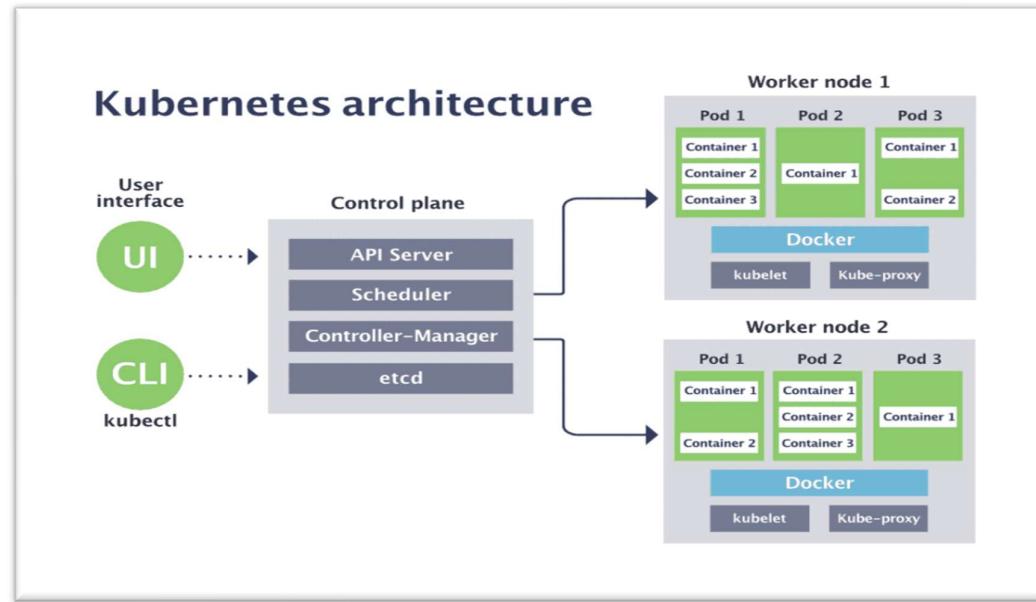
4. Worker Node

- Executes the containers.
- Managed by the master node.
- **Components:**
 - [kubelet](#): Communicates with the control plane
 - [kube-proxy](#): Network communication
 - [Container Runtime](#) (e.g., containerd or Docker)

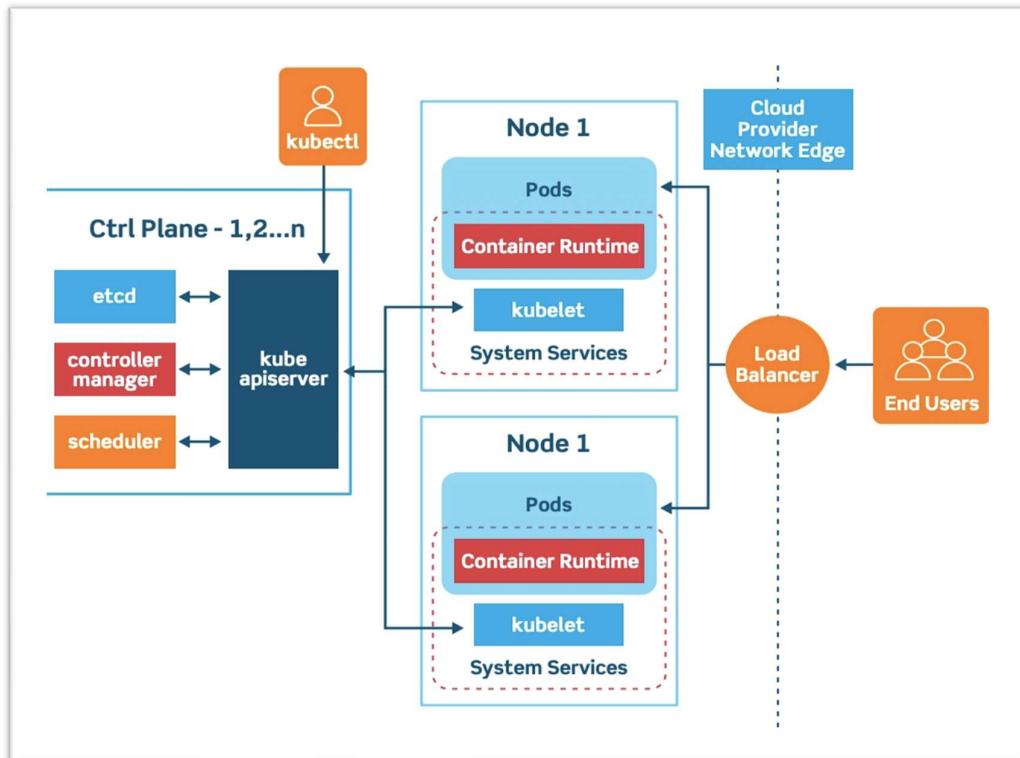
5. Pods

- The smallest deployable unit in Kubernetes.
- A pod contains one or more containers with shared storage/network.
- Pods are ephemeral (short-lived).

Kubernetes Architecture



Working:



Azure Kubernetes Service (AKS)

- AKS is a managed Kubernetes service in Microsoft Azure.
- Azure handles:
 - Control plane (master nodes)
 - Upgrades
 - Monitoring
- You manage:
 - Worker nodes
 - Deployments
 - Networking

Benefits of AKS:

- Simplified deployment
- Integrated with Azure Monitor and AAD
- Built-in scaling & self-healing

kubectl – Kubernetes CLI Tool

Used to interact with Kubernetes clusters.

Basic Commands

Command	Description
kubectl get nodes	List all nodes in the cluster
kubectl get pods	List all running pods
Kubectl get pod --namespace=<name>	List all running pods in the given namespace
Kubectl create namespace <name>	Creates a namespace with given name

Command	Description
Kubectl run <pod-name> --image=<img-name>	Creates a pod for an image in docker hub's public repository
kubectl get services	List all services
Kubectl get namespaces	List all namespaces in the cluster
kubectl apply -f <file.yaml>	Apply configuration from a YAML file
kubectl create -f <file.yaml>	Create resources from YAML
kubectl delete -f <file.yaml>	Delete resources from YAML
kubectl describe pod <pod-name>	Detailed info about a pod
kubectl logs <pod-name>	View logs from a pod
kubectl exec -it <pod-name> -- /bin/bash	Access a running container inside a pod

AKS Specific Commands

1. Login to Azure CLI

[az login](#)

2. Get Credentials for Your AKS Cluster

[az aks get-credentials --resource-group <resource-group-name> --name <aks-cluster-name>](#)

3. Check Cluster Context

[kubectl config get-contexts](#)

[kubectl config use-context <context-name>](#)

Kubernetes Manifest Files

Kubernetes uses **YAML manifest files** to define the desired state of cluster resources. These files declare what you want (e.g., a pod, a service, a deployment), and Kubernetes takes care of making it happen.

Each YAML file typically includes:

- apiVersion
- kind (resource type - object like pod)
- metadata (name, labels)
- spec (specifications- containers,ports)

1. Pod Manifest File

A **Pod** is the smallest deployable unit that runs one or more containers.

Sample pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: my-nginx-pod
labels:
  app: nginx
spec:
  containers:
    - name: nginx-container
      image: nginx
      ports:
        - containerPort: 80
```

Command to Apply:

```
kubectl apply -f pod.yaml
```

2. Service Manifest File

A **Service** exposes your pod(s) to other services or external users. Common types:

- ClusterIP (internal)
- NodePort (external on node port)
- LoadBalancer (for cloud-based external access)

Sample service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  type: NodePort
  ports:
    - port: 80
      targetPort: 80
      nodePort: 30007
```

Command to Apply:

```
kubectl apply -f service.yaml
```

3. Deployment Manifest File

A **Deployment** manages pods by maintaining the desired number of replicas, rolling updates, etc.

Sample deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
```

```
name: nginx-deployment

spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
```

 **Command to Apply:**

```
kubectl apply -f deployment.yaml
```

Working Process: Master Node to Worker Node in Kubernetes

1. User Applies Manifest

```
kubectl apply -f deployment.yaml
```

- This command goes to the **API Server** on the **Master Node**.

2. API Server

- Validates the request.
- Updates the **etcd** database with the desired cluster state.

3. Scheduler

- Checks which **Worker Node** has available resources (CPU, memory).

- Chooses the best node and schedules the **pod** there.

4. Controller Manager

- Ensures the desired state is achieved.
- It sees: "We want 3 pods of NGINX" → ensures they are always running.
- If any pod fails, it reschedules them.

5. Kubelet (on Worker Node)

- Master node instructs **Kubelet** (agent on the Worker).
- Kubelet receives pod specs and creates the container via the **container runtime** (Docker, containerd).

6. Container Runtime

- Pulls the image (e.g., nginx) from a container registry (like Docker Hub).
- Runs the container inside a pod.

7. kube-proxy

- Manages networking rules to route requests to the correct pod.
- Handles load balancing for services.



Ongoing Communication

- **Health Checks:** Kubelet sends regular status reports to the Master.
- **Events:** If a pod dies, Master is notified → Controller reschedules it.
- **Scaling:** User changes replicas in deployment.yaml → Scheduler creates or deletes pods accordingly.

Some kubectl commands:

```
gouth@Priya_gowtami MINGW64 ~
$ kubectl get namespace
NAME        STATUS   AGE
default     Active  2d17h
gani-kube   Active  22h
jyoshna-meghadri  Active  17h
kube-node-lease  Active  2d17h
kube-public   Active  2d17h
kube-system   Active  2d17h
madhusrivamisett  Active  15h
poc-o1       Active  4h25m
priya-ns     Active  15h
reshma-node   Active  17h
testns      Active  38h

gouth@Priya_gowtami MINGW64 ~
$ kubectl create namespace example-ns
namespace/example-ns created
```

```
gouth@Priya_gowtami MINGW64 ~
$ kubectl get namespace
NAME        STATUS   AGE
default     Active  2d17h
example-ns  Active  5s
gani-kube   Active  22h
jyoshna-meghadri  Active  17h
kube-node-lease  Active  2d17h
kube-public   Active  2d17h
kube-system   Active  2d17h
madhusrivamisett  Active  15h
poc-o1       Active  4h27m
priya-ns     Active  15h
reshma-node   Active  17h
testns      Active  38h
```

```
gouth@Priya_gowtami MINGW64 ~
$ kubectl run nginx --image=nginx --namespace=example-ns
pod/nginx created
```

```
gouth@Priya_gowtami MINGW64 ~
$ kubectl get pod --namespace=example-ns
NAME    READY  STATUS    RESTARTS   AGE
nginx  1/1    Running   0          16s

gouth@Priya_gowtami MINGW64 ~
$ kubectl get service
NAME        TYPE           CLUSTER-IP    EXTERNAL-IP
kubernetes  ClusterIP     10.0.0.1     <none>
nginx-svc   LoadBalancer   10.0.14.236  135.235.248.95
```

```
gouth@Priya_gowtami MINGW64 ~
$ kubectl delete pod nginx --namespace=example-ns
pod "nginx" deleted
gouth@Priya_gowtami MINGW64 ~
$ kubectl get pod --namespace=example-ns
No resources found in example-ns namespace.
```

```
gouth@Priya_gowtami MINGW64 ~
$ kubectl delete namespace example-ns
namespace "example-ns" deleted
```

```
gouth@Priya_gowtami MINGW64 ~
$ kubectl get namespace
NAME        STATUS   AGE
default     Active  2d17h
gani-kube   Active  22h
jyoshna-meghadri  Active  17h
kube-node-lease  Active  2d17h
kube-public   Active  2d17h
kube-system   Active  2d17h
madhusrivamisett  Active  15h
poc-o1       Active  4h34m
priya-ns     Active  16h
reshma-node   Active  17h
testns      Active  38h
```